

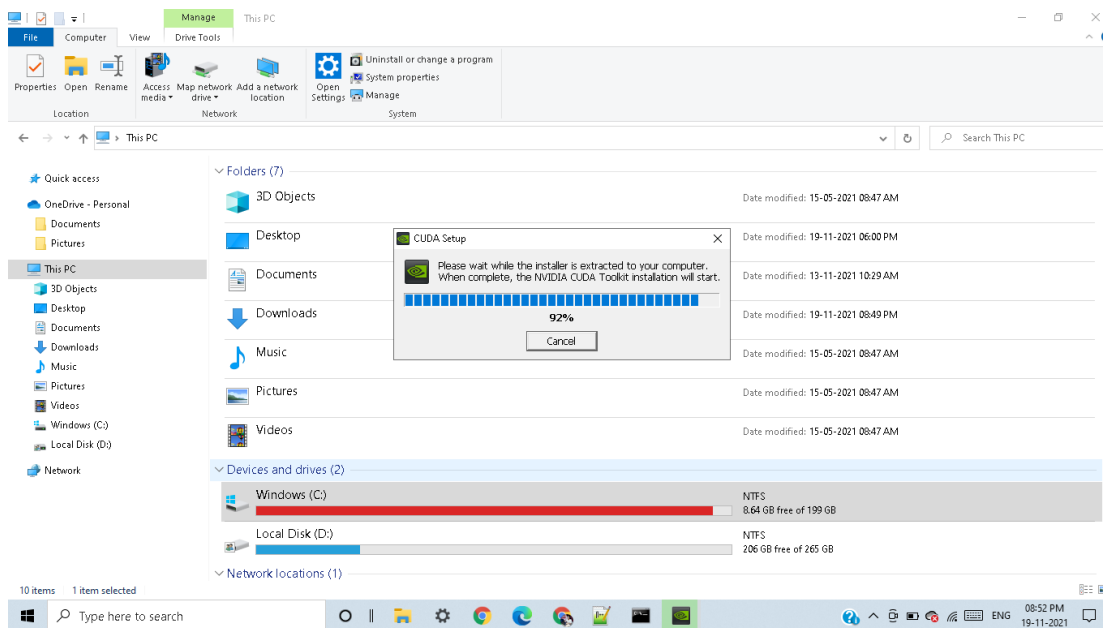
## **Practical No. 7**

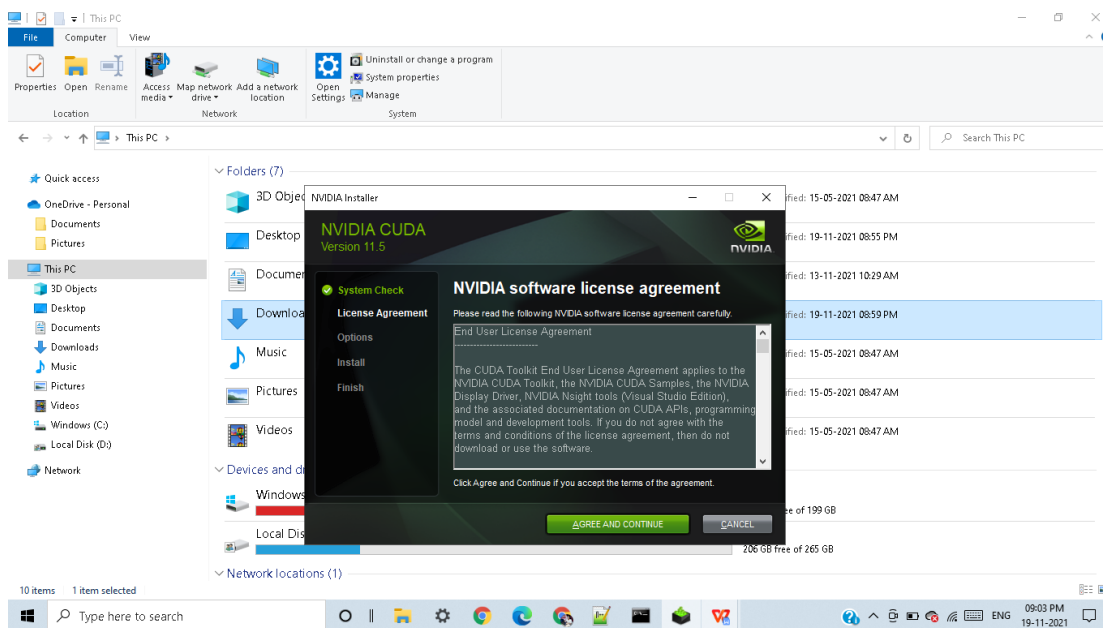
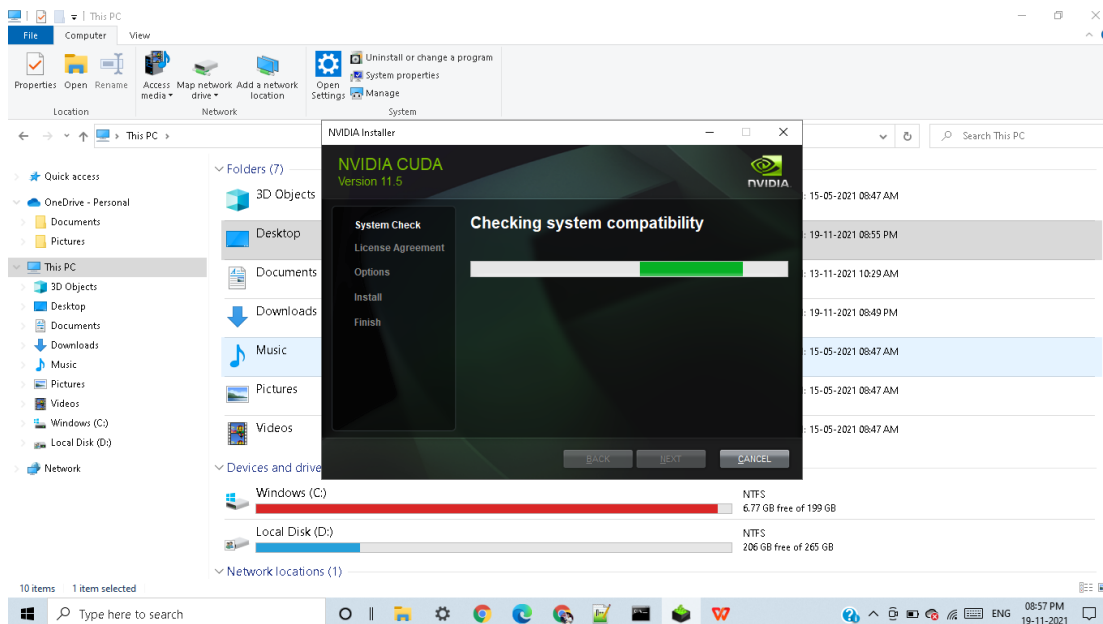
### **Study and Implementation of CUDA C basic functions.**



**NAME: SHWETA NANDKUMAR ARBUNE**  
**PRN: 2019BTECS00205**

- 1) Setup the environment requirements, for execution of CUDA C programs.


#### **A. CUDA Toolkit Installation :**



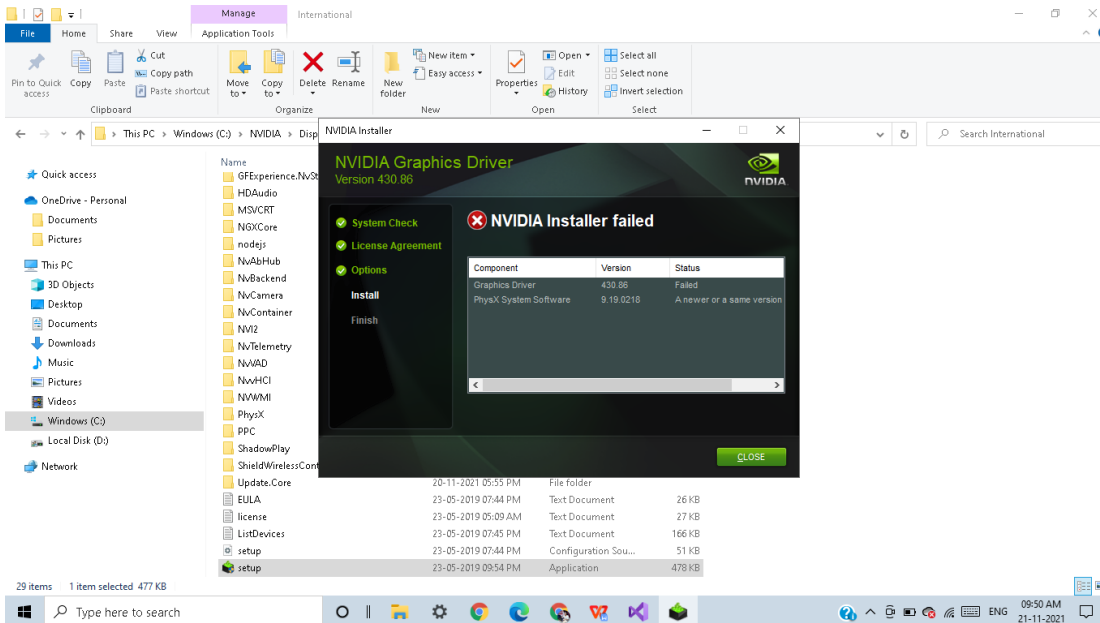




## B. Showing Error of CudaDrivers:



### C. CUDA Drivers Installation Failed :



### D. CUDA Setup on Google colab:

Cuda1.ipynb - Colaboratory

colab.research.google.com/drive/1uzH-lYDQxJm\_d6sn0UX7z\_pWReNNN

Cuda1.ipynb

File Edit View Insert Runtime Tools Help

+ Code + Text

```
lapt-get --purge remove cuda nvidia* libnvidia-*
dpkg -l | grep cuda- | awk '{print $2}' | xargs -n1 dpkg --purge
lapt-get remove cuda-*
lapt autoremove
lapt-get update
```

Reading package lists... Done  
Building dependency tree  
Reading state information... Done  
Note, selecting 'nvidia-kernel-common-418-server' for glob 'nvidia\*'  
Note, selecting 'nvidia-325-updates' for glob 'nvidia\*'  
Note, selecting 'nvidia-346-updates' for glob 'nvidia\*'  
Note, selecting 'nvidia-driver-binary' for glob 'nvidia\*'  
Note, selecting 'nvidia-331-dev' for glob 'nvidia\*'  
Note, selecting 'nvidia-394-updates-dev' for glob 'nvidia\*'  
Note, selecting 'nvidia-compute-utils-418-server' for glob 'nvidia\*'  
Note, selecting 'nvidia-384-dev' for glob 'nvidia\*'  
Note, selecting 'nvidia-libopencl-346-updates' for glob 'nvidia\*'  
Note, selecting 'nvidia-fs-prebuilt' for glob 'nvidia\*'  
Note, selecting 'nvidia-driver-440-server' for glob 'nvidia\*'  
Note, selecting 'nvidia-340-updates-uvn' for glob 'nvidia\*'  
Note, selecting 'nvidia-dkms-450-server' for glob 'nvidia\*'  
Note, selecting 'nvidia-kernel-common' for glob 'nvidia\*'  
Note, selecting 'nvidia-kernel-source-440-server' for glob 'nvidia\*'  
Note, selecting 'nvidia-gds' for glob 'nvidia\*'  
Note, selecting 'nvidia-331-updates-uvn' for glob 'nvidia\*'  
Note, selecting 'nvidia-cg-toolkit' for glob 'nvidia\*'

Cuda1.ipynb - Colaboratory

colab.research.google.com/drive/1uzH-lYDQxJm\_d6sn0UX7z\_pWReNNN

Cuda1.ipynb

File Edit View Insert Runtime Tools Help

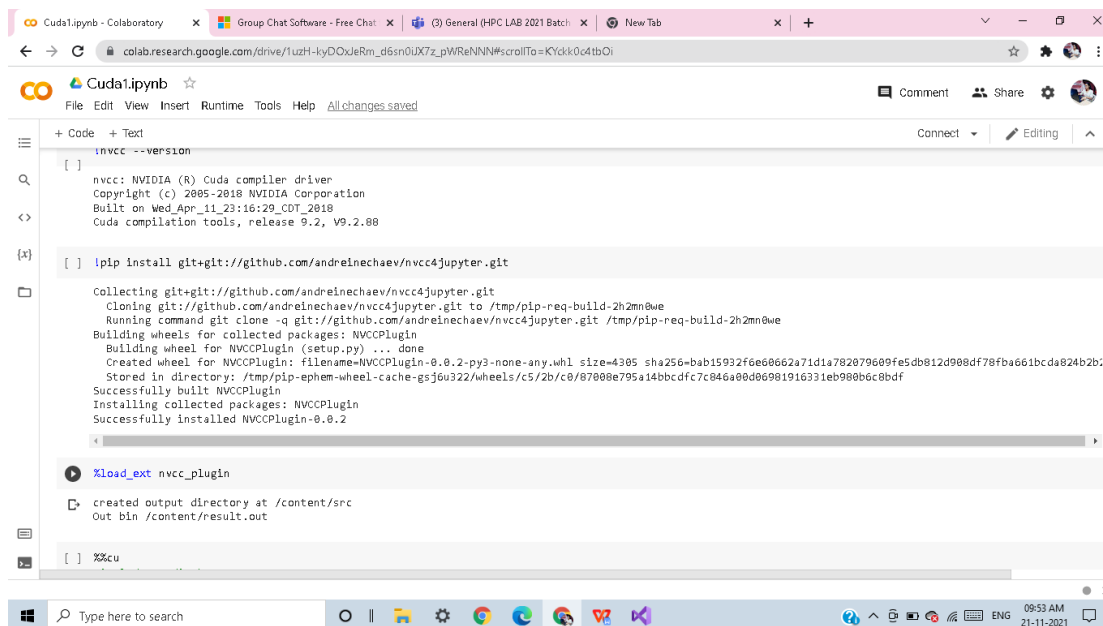
+ Code + Text

```
Note, selecting 'nvidia-prime' for glob 'nvidia*'  
Note, selecting 'nvidia-dkms-390' for glob 'nvidia*'  
Note, selecting 'nvidia-kernel-dkms' for glob 'nvidia*'
```

```
!wget https://developer.nvidia.com/compute/cuda/9.2/Prod/local_installers/cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64  
dpkg -i cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb  
lapt-key add /var/cuda-repo-9-2-local/7fa2af80.pub  
lapt-get update  
lapt-get install cuda-9.2
```

```
--2021-11-20 04:59:13-- https://developer.nvidia.com/compute/cuda/9.2/Prod/local_installers/cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64  
Resolving developer.nvidia.com (developer.nvidia.com)... 152.195.19.142  
Connecting to developer.nvidia.com (developer.nvidia.com)|152.195.19.142|:443... connected.  
HTTP request sent, awaiting response... 301 Moved Permanently  
Location: https://developer.nvidia.com/compute/cuda/9.2/Prod/local_installers/cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64 [following]  
--2021-11-20 04:59:13-- https://developer.nvidia.com/compute/cuda/9.2/Prod/local_installers/cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64  
Reusing existing connection to developer.nvidia.com:443.  
HTTP request sent, awaiting response... 302 Found  
Location: https://developer.download.nvidia.com/compute/cuda/9.2/secure/Prod/local_installers/cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb?fyGsp7_89  
--2021-11-20 04:59:14-- https://developer.download.nvidia.com/compute/cuda/9.2/secure/Prod/local_installers/cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64  
Resolving developer.download.nvidia.com (developer.download.nvidia.com)... 152.195.19.142  
Connecting to developer.download.nvidia.com (developer.download.nvidia.com)|152.195.19.142|:443... connected.  
HTTP request sent, awaiting response... 200 OK  
Length: 1267391958 (1.26) [application/x-deb]  
Saving to: 'cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb'  
  
cuda-repo-ubuntu160 100%[=====] 1.18G 150MB/s in 7.9s  
  
2021-11-20 04:59:22 (153 MB/s) - 'cuda-repo-ubuntu1604-9-2-local_9.2.88-1_amd64.deb' saved [1267391958/1267391958]
```

https://developer.download.nvidia.com/compute/cuda/9.2/secure/Prod/local\_installers/cuda-repo-ubuntu1604-9-2-local\_9.2.88-1\_amd64.deb?fyGsp7\_89&P7\_89&krak876NT-EksAF9T87DnuCyF5K10itdE3XRPTNjgtnX-O20-GvXp\_Oehcyg4fobo-aoovyZVw1...



```
nvcc --version
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2018 NVIDIA Corporation
Built on Wed_Apr_11_23:16:29_CDT_2018
Cuda compilation tools, release 9.2, V9.2.88

[ ] !pip install git+git://github.com/andreinechaev/nvcc4jupyter.git

Collecting git+git://github.com/andreinechaev/nvcc4jupyter.git
  Cloning git://github.com/andreinechaev/nvcc4jupyter.git to /tmp/pip-req-build-2h2mn0we
  Running command git clone -q git://github.com/andreinechaev/nvcc4jupyter.git /tmp/pip-req-build-2h2mn0we
  Building wheels for collected packages: NVCCPlugin
  Building wheel for NVCCPlugin (setup.py) ... done
  Created wheel for NVCCPlugin: filename=NVCCPlugin-0.0.2-py3-none-any.whl size=4305 sha256=bab15932f6ee0662a71d1a782079609fe5db812d908df78fba661bcd824b2b1
  Stored in directory: /tmp/pip-ephem-wheel-cache-gs36u322/wheels/c5/2b/c0/870808e795a14bbcdfc7c046a08d06901916331eb980b6c8bdf
  Successfully built NVCCPlugin
  Installing collected packages: NVCCPlugin
  Successfully installed NVCCPlugin-0.0.2

%load_ext nvcc_plugin
created output directory at /content/src
Out bin /content/result.out

[ ] %%cu
```

2) Execute the attached Program 1, and understand the output.

There is 1 device supporting CUDA

```
Device 0: "Tesla K80"
Major revision number:      3
Minor revision number:      7
Total amount of global memory:  -887947264 bytes
Total amount of constant memory: 65536 bytes
Total amount of shared memory per block: 49152 bytes
Total number of registers available per block: 65536
Warp size: 32
Multiprocessor count: 13
Maximum number of threads per block: 1024
Maximum sizes of each dimension of a block: 1024 x 1024 x 64
Maximum sizes of each dimension of a grid: 2147483647 x 65535 x 65535
Maximum memory pitch: 2147483647 bytes
Texture alignment: 512 bytes
Clock rate: 562000 kilohertz
```

✓ 0s completed at 11:26 AM

Device\_Count is basically used to check the GPU device available on system or not, In google colab here 1 device with id 0 is present and the name of device is Tesla k80. Major and minor revision number is printed after it. Global memory is checked with deviceProp.totalGlobalMem function in bytes same constant memory and shared memory is checked for SM. Multiprocessor cont and threads per block is printed after it.

- 3) Write a CUDA C program to perform the addition of two vectors of arbitrary size (Dynamic Array).

CODE:

```
%%cu
#include <stdint.h>
#include<stdio.h>
#include<unistd.h>

__global__ void vectoradd(int *x,int *y, int *z)
{
    int id=blockIdx.x;
    z[id]=x[id]+y[id];
}

int main(void)
{
    int a[8];
    int b[8];
    int c[8];
    int *d,*e,*f;
    int i;
    for(i=0;i<8;i++)
    {
        a[i]=i;
        b[i]=i;
    }

    cudaMalloc((void **)&d,8*sizeof(int));
    cudaMalloc((void **)&e,8*sizeof(int));
    cudaMalloc((void **)&f,8*sizeof(int));

    cudaMemcpy(d,a,8*sizeof(int),cudaMemcpyHostToDevice);
    cudaMemcpy(e,b,8*sizeof(int),cudaMemcpyHostToDevice);

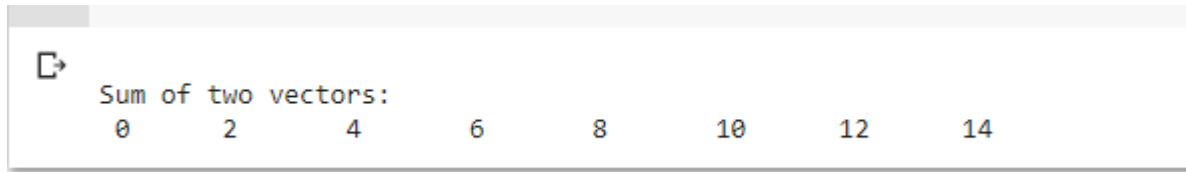
    vectoradd<<<8,1>>>>(d,e,f);

    cudaMemcpy(c,f,8*sizeof(int),cudaMemcpyDeviceToHost);

    printf("\nSum of two vectors:\n ");
    for(i=0;i<8;i++)
    {
        printf("%d\t",c[i]);
    }
}
```

```
    cudaFree(d);  
    cudaFree(e);  
    cudaFree(f);  
  
    return 0;  
}
```

OUTPUT:

A terminal window with a light gray header bar. The output text is as follows:

```
Sum of two vectors:  
0      2      4      6      8     10     12     14
```

Github link:

<https://github.com/shwetaarbune/HPC-LAB7>